

Joel Schwarzmann

Product & Engineering Leadership / Data Platforms / Developer Tooling

📍 London, England, GB ✉ joel.l.schw@gmail.com 🔗 [datajoely](#)

Builder of developer tooling and data / AI platforms. Combines hands-on engineering with product leadership to deliver secure, interoperable systems. Skilled at translating complex technical concepts for diverse audiences. Passionate about creating platforms and products users love.

Experience

Director of Platform

2024-09-01 - Present

Aneira Health

Leading development of a GCP-native data platform providing women with personalised, clinically validated healthcare insights that directly improve patient outcomes. Simultaneously building a world-class multimodal research platform integrating de-identified EHR, genomics, DICOM, and wearable data.

- Delivered a production grade data platform integrating EHR, genomics, and wearable data streams.
- Established FHIR-compliant architecture with interoperability as a core feature, including an MCP layer to enable secure LLM use cases (e.g., triage, coaching).
- Directed product discovery and initial development to expose the research platform to the first external pharmaceutical customer.
- Operated in a seed-stage context, spanning product, platform, and engineering execution.

Advisor

2024-09-01 - Present

Population Health Partners (incubator of Aneira Health among other companies)

Strategic and technical advisor across portfolio companies on data, ML, and AI platforms.

- Guided ML/AI capability development, engineering standards, and scalable data infrastructure across biotech, pharmaceutical, and healthcare portfolio companies.
- Led vendor assessments and selections (e.g., Vanta, Dagster, Linear, Miro, Posthog).

Principal Technical Product Manager

2021-01-01 - 2024-08-31

QuantumBlack Labs (a McKinsey company)

Product leader for developer tooling and durable ML delivery; stewarded open source and incubated multiple products end-to-end.

- Incubated an internal multi-tenant 'ML Models as a Service' platform from Aug 2023 to initial client deployments in 2024.
- Created Alloy, an internal build tool to compose Kedro projects from reusable building blocks; simplified packaging, CI/CD, and testing, driving organisation-wide adoption.
- Drove velocity on the open-source project Kedro (10.5k ★) first as a contributor and then as a product lead, leading efforts across product research, developer relations, and developing overall product toolkit. Stewarded Kedro through the Linux Foundation for AI & Data from incubation to graduated project. Drove a customer-driven research initiative that overhauled Kedro's dynamic configuration experience, significantly reducing overhead. Continue to serve on the Kedro Technical Steering Committee (TSC) outside of the organisation.

- Directed cross-functional teams to drive consistency and raise engineering excellence across a global ML cross-industry product portfolio.
- Defined privacy-preserving product telemetry standards and rollout practices.
- Organised developer community events and Demo Days to drive reuse and collaboration.

Senior Data Engineer

2017-02-28 - 2020-12-31

QuantumBlack (a McKinsey company)

Forward-deployed Data Engineer delivering production ML across pharma, manufacturing, insurance, and motorsport.

- Built reinforcement-learning-driven race strategy software for a championship-winning Formula E team; led Phase 1 (Engineering) and Phase 2 (Product).
- Designed and implemented feature engineering and MLOps deployment pipelines across vaccine production, diabetes patient journeys, motor insurance pricing, water treatment amongst others.

Senior Forensic Data Analyst

2013-09-01 - 2017-02-17

PwC

Delivered data-driven investigations into fraud, disputes, and trade surveillance; developed SQL/Python toolkits.

- Analysed Bloomberg communications related to LIBOR front-running.
- Resolved creditor disputes and managed decommissioning projects following the Lehman Brothers bankruptcy.
- Led development of graph-database-powered market surveillance tooling.

Education

BSc in Computer Science with Management

2010-09-01 - 2013-07-01

University of Birmingham

Grade: First Class (Hons)

Awards

Best Final Year Student

University of Birmingham, School of Computer Science

Skills

Languages: Python, SQL

Frameworks: Dagster, Airflow, dbt, sqlmesh, PySpark, Kedro, Ibis, Kafka, Pandera, FastMCP, FastAPI

Product & Practices: Product Management, Platform Strategy, MLOps, Telemetry & Privacy, DevEx, CI/CD, LLMs, Testing

Projects

Kedro (Linux Foundation for AI & Data)

Kedro is a toolbox for production-ready data science. It uses software engineering best practices to help create data engineering and data science pipelines that are reproducible, maintainable, and modular.

Reasoning Over Pre-training: Evaluating LLM Performance and Augmentation in Women's Health	2025-05-23
<i>medRxiv</i>	